

Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

```
In [2]: df=pd.read_csv("madrid_2002.csv")
df
```

Out[2]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10
0	2002-04-01 01:00:00	NaN	1.39	NaN	NaN	NaN	145.100006	352.100006	NaN	6.54	41.990002
1	2002-04-01 01:00:00	1.93	0.71	2.33	6.20	0.15	98.150002	153.399994	2.67	6.85	20.980000
2	2002-04-01 01:00:00	NaN	0.80	NaN	NaN	NaN	103.699997	134.000000	NaN	13.01	28.440001
3	2002-04-01 01:00:00	NaN	1.61	NaN	NaN	NaN	97.599998	268.000000	NaN	5.12	42.180000
4	2002-04-01 01:00:00	NaN	1.90	NaN	NaN	NaN	92.089996	237.199997	NaN	7.28	76.330002
...
217291	2002-11-01 00:00:00	4.16	1.14	NaN	NaN	NaN	81.080002	265.700012	NaN	7.21	36.750000
217292	2002-11-01 00:00:00	3.67	1.73	2.89	NaN	0.38	113.900002	373.100006	NaN	5.66	63.389999
217293	2002-11-01 00:00:00	1.37	0.58	1.17	2.37	0.15	65.389999	107.699997	1.30	9.11	9.640000
217294	2002-11-01 00:00:00	4.51	0.91	4.83	10.99	NaN	149.800003	202.199997	1.00	5.75	NaN
217295	2002-11-01 00:00:00	3.11	1.17	3.00	7.77	0.26	80.110001	180.300003	2.25	7.38	29.240000

217296 rows × 16 columns

Data Cleaning and Data Preprocessing

```
In [3]: df=df.dropna()
```

```
In [4]: df.columns
```

```
Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 32381 entries, 1 to 217295
Data columns (total 16 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      32381 non-null   object 
 1   BEN        32381 non-null   float64
 2   CO         32381 non-null   float64
 3   EBE        32381 non-null   float64
 4   MXY        32381 non-null   float64
 5   NMHC       32381 non-null   float64
 6   NO_2       32381 non-null   float64
 7   NOx        32381 non-null   float64
 8   OXY        32381 non-null   float64
 9   O_3         32381 non-null   float64
 10  PM10       32381 non-null   float64
 11  PXY        32381 non-null   float64
 12  SO_2       32381 non-null   float64
 13  TCH        32381 non-null   float64
 14  TOL        32381 non-null   float64
 15  station    32381 non-null   int64  
dtypes: float64(14), int64(1), object(1)
memory usage: 4.2+ MB
```

```
In [6]: data=df[['CO' , 'station']]  
data
```

Out[6]:

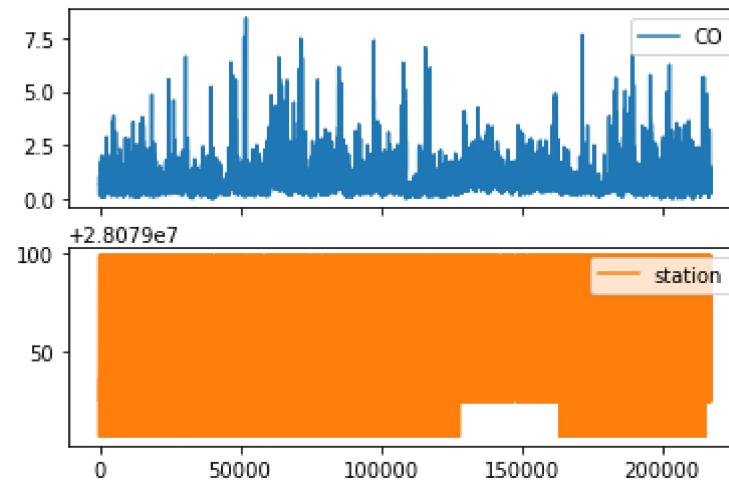
	CO	station
1	0.71	28079035
5	0.72	28079006
22	0.80	28079024
24	1.04	28079099
26	0.53	28079035
...
217269	0.28	28079024
217271	1.30	28079099
217273	0.97	28079035
217293	0.58	28079024
217295	1.17	28079099

32381 rows × 2 columns

Line chart

```
In [7]: data.plot.line(subplots=True)
```

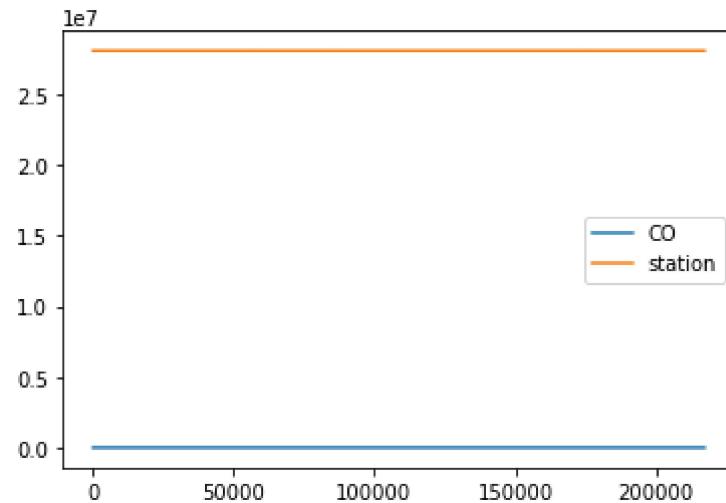
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)



Line chart

```
In [8]: data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```

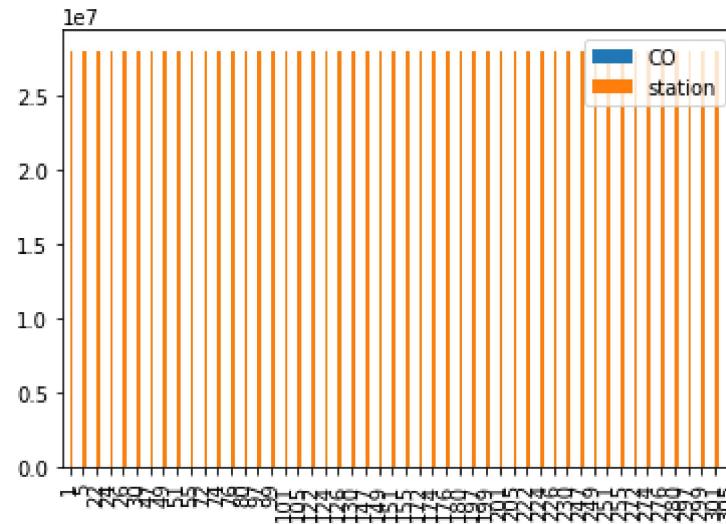


Bar chart

```
In [9]: b=data[0:50]
```

```
In [10]: b.plot.bar()
```

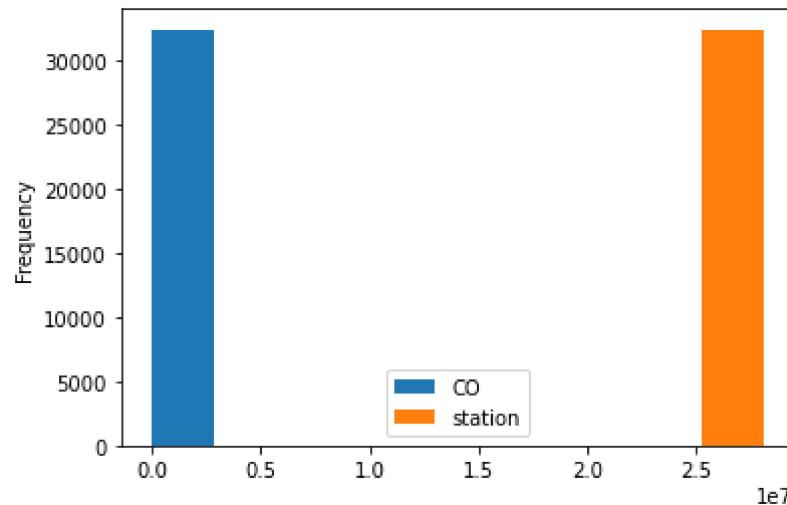
```
Out[10]: <AxesSubplot:>
```



Histogram

```
In [11]: data.plot.hist()
```

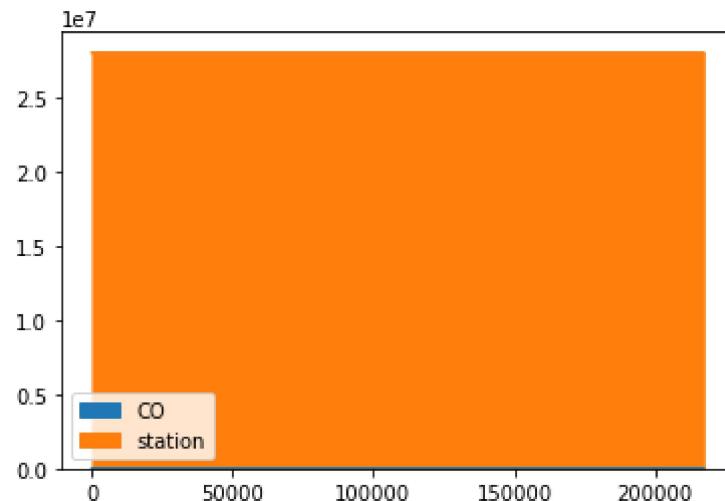
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

```
In [12]: data.plot.area()
```

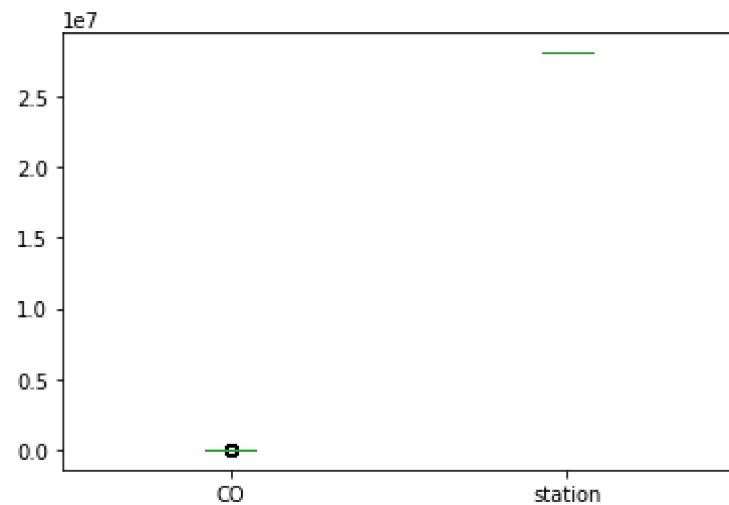
```
Out[12]: <AxesSubplot:>
```



Box chart

In [13]: `data.plot.box()`

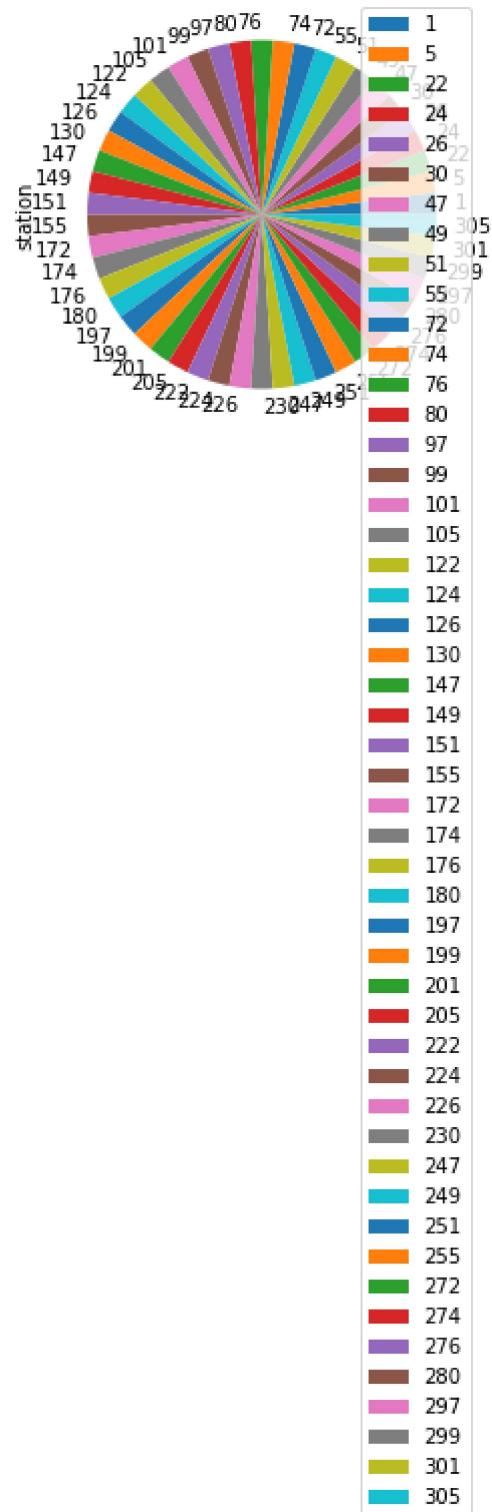
Out[13]: <AxesSubplot:>



Pie chart

```
In [14]: b.plot.pie(y='station' )
```

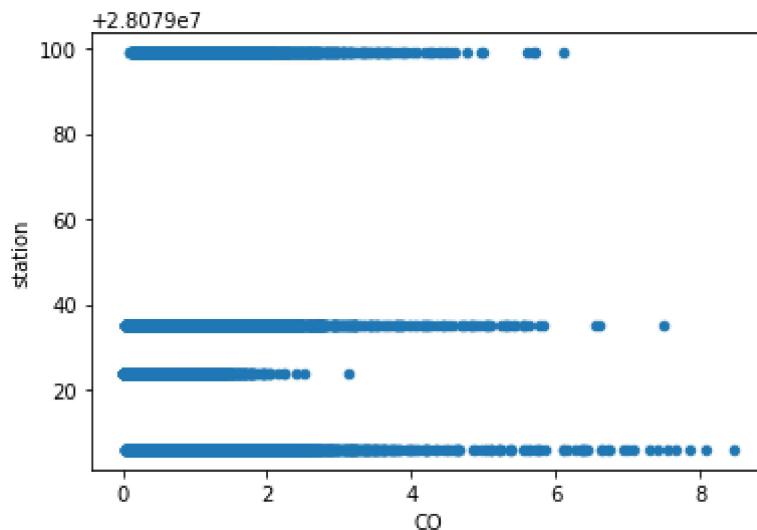
```
Out[14]: <AxesSubplot:ylabel='station'>
```



Scatter chart

```
In [15]: data.plot.scatter(x='CO' ,y='station')
```

```
Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [16]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 32381 entries, 1 to 217295
Data columns (total 16 columns):
 #   Column   Non-Null Count  Dtype  
--- 
 0   date      32381 non-null   object 
 1   BEN       32381 non-null   float64
 2   CO        32381 non-null   float64
 3   EBE       32381 non-null   float64
 4   MXY       32381 non-null   float64
 5   NMHC      32381 non-null   float64
 6   NO_2      32381 non-null   float64
 7   NOx       32381 non-null   float64
 8   OXY       32381 non-null   float64
 9   O_3        32381 non-null   float64
 10  PM10      32381 non-null   float64
 11  PXY       32381 non-null   float64
 12  SO_2      32381 non-null   float64
 13  TCH       32381 non-null   float64
 14  TOI       32381 non-null   float64
```

In [17]: df.describe()

Out[17]:

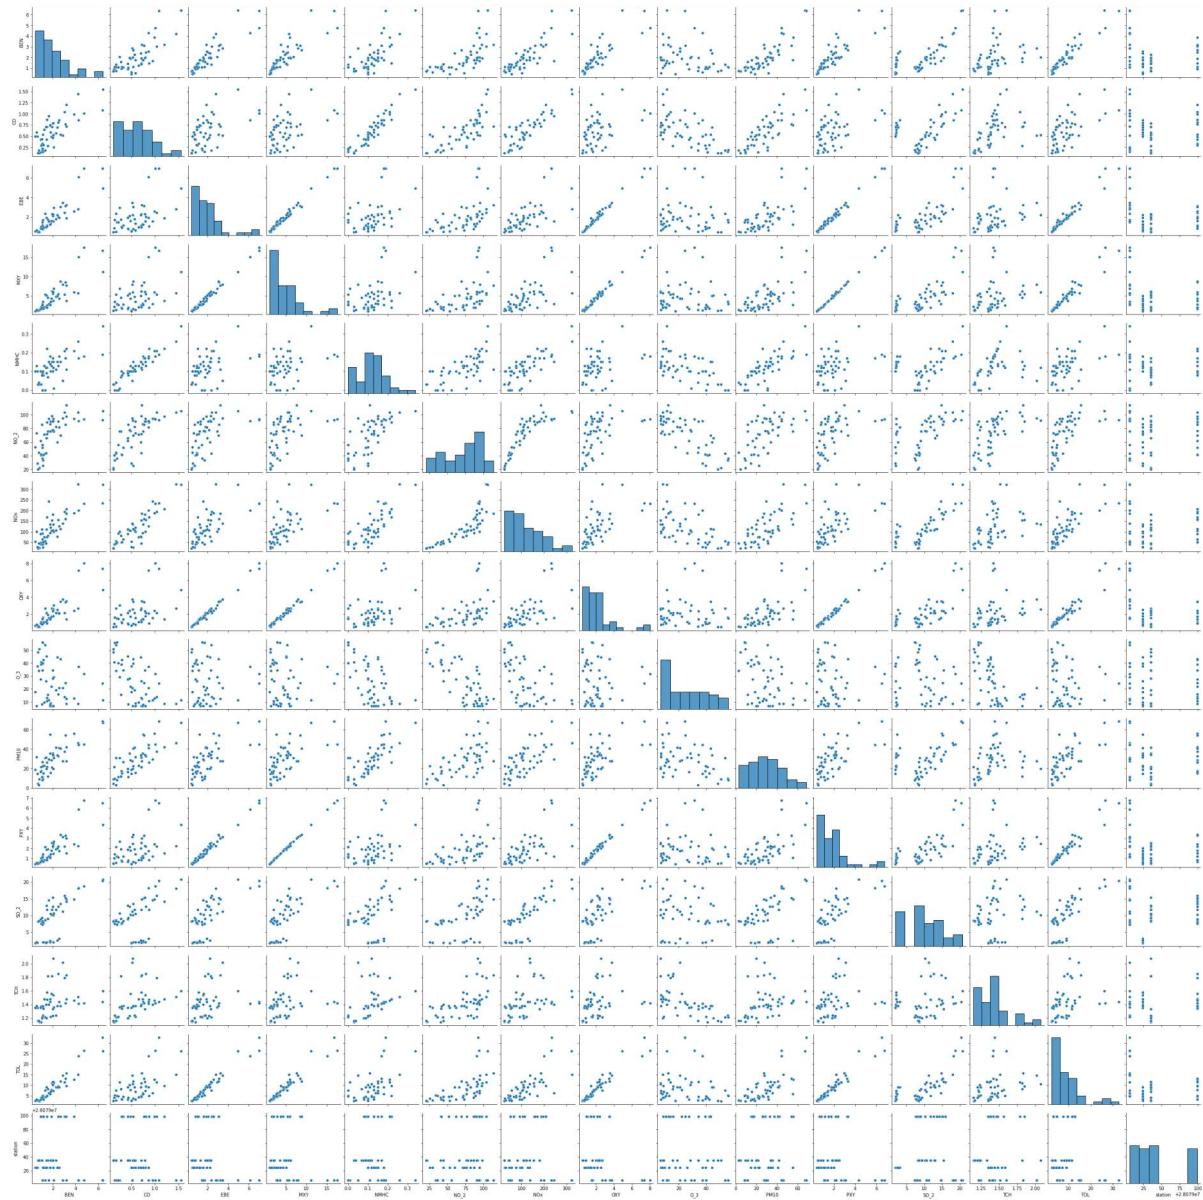
	BEN	CO	EBE	MXY	NMHC	NO_2	
count	32381.000000	32381.000000	32381.000000	32381.000000	32381.000000	32381.000000	323
mean	2.479155	0.787323	2.914004	7.013636	0.155827	58.936796	1
std	2.280959	0.610810	2.667881	6.774365	0.135731	31.472733	1
min	0.180000	0.000000	0.180000	0.190000	0.000000	0.890000	
25%	0.970000	0.420000	1.140000	2.420000	0.080000	35.660000	
50%	1.840000	0.620000	2.130000	5.140000	0.130000	57.160000	
75%	3.250000	0.980000	3.830000	9.420000	0.200000	78.769997	1
max	32.660000	8.460000	41.740002	99.879997	2.700000	263.600006	13

In [18]: df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]

EDA AND VISUALIZATION

```
In [19]: sns.pairplot(df1[0:50])
```

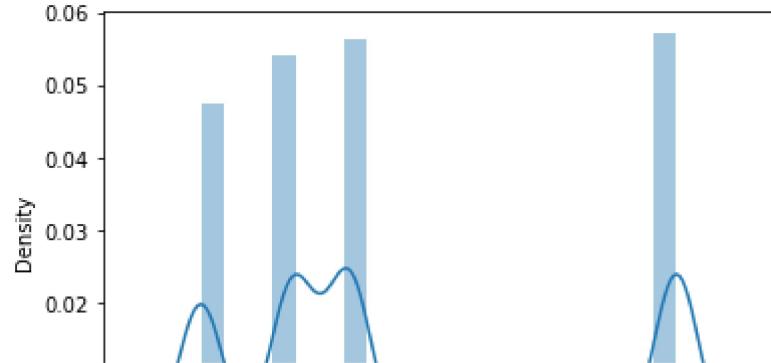
```
Out[19]: <seaborn.axisgrid.PairGrid at 0x1c1d56f0460>
```



In [20]: `sns.distplot(df1['station'])`

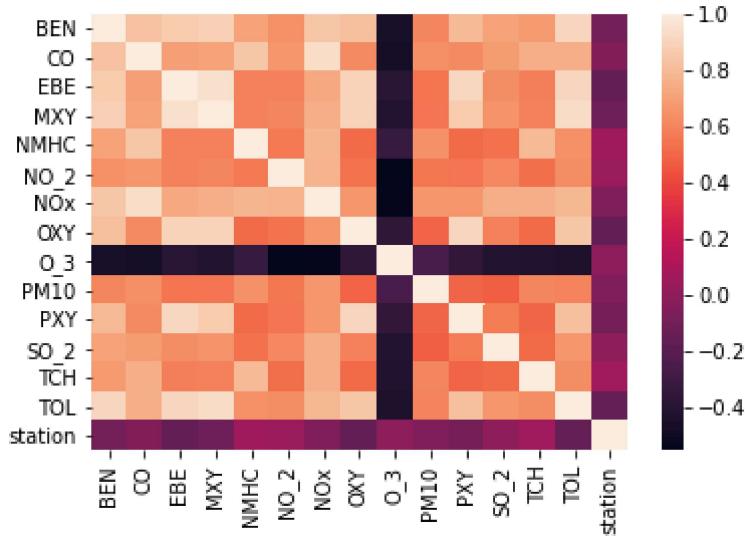
```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)
```

Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>



In [21]: `sns.heatmap(df1.corr())`

Out[21]: <AxesSubplot:>



TO TRAIN THE MODEL AND MODEL BUILDING

In [22]: `x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y=df['station']`

```
In [23]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: lr.intercept_
```

```
Out[25]: 28078987.901386347
```

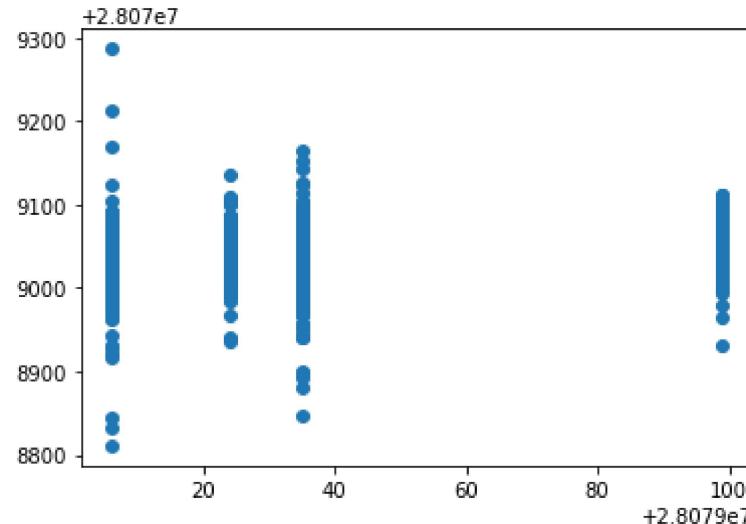
```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

```
Out[26]:
```

	Co-efficient
BEN	2.125454
CO	-13.766248
EBE	-11.419364
MXY	4.140084
NMHC	83.042808
NO_2	0.255394
NOx	-0.098849
OXY	-5.100131
O_3	-0.033571
PM10	-0.115078
PXY	7.991349
SO_2	0.671541
TCH	43.578380
TOL	-1.661425

```
In [27]: prediction = lr.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x1c1e4f26580>
```



ACCURACY

```
In [28]: lr.score(x_test,y_test)
```

```
Out[28]: 0.1837216607188462
```

```
In [29]: lr.score(x_train,y_train)
```

```
Out[29]: 0.20410773820018646
```

Ridge and Lasso

```
In [30]: from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

Accuracy(Ridge)

```
In [32]: rr.score(x_test,y_test)
```

```
Out[32]: 0.18324514844316653
```

```
In [33]: rr.score(x_train,y_train)
```

```
Out[33]: 0.20390330994003047
```

```
In [34]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

```
In [35]: la.score(x_train,y_train)
```

```
Out[35]: 0.061940553181437985
```

Accuracy(Lasso)

```
In [36]: la.score(x_test,y_test)
```

```
Out[36]: 0.05426985291531594
```

```
In [37]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: en.coef_
```

```
Out[38]: array([ 1.11180324,  0.          , -2.79863616,  1.53404939,  0.23062826,
   0.22914989, -0.02762068, -2.37311702, -0.02536318,  0.01281824,
   2.32523214,  0.47914468,  1.11459539, -1.30459123])
```

```
In [39]: en.intercept_
```

```
Out[39]: 28079038.287570737
```

```
In [40]: prediction=en.predict(x_test)
```

```
In [41]: en.score(x_test,y_test)
```

```
Out[41]: 0.08795788332481069
```

Evaluation Metrics

```
In [42]: from sklearn import metrics  
print(metrics.mean_absolute_error(y_test,prediction))  
print(metrics.mean_squared_error(y_test,prediction))  
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
28.481753575637345  
1109.382765069125  
33.30739805312215
```

Logistic Regression

```
In [43]: from sklearn.linear_model import LogisticRegression
```

```
In [44]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_P10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
target_vector=df['station']
```

```
In [45]: feature_matrix.shape
```

```
Out[45]: (32381, 14)
```

```
In [46]: target_vector.shape
```

```
Out[46]: (32381,)
```

```
In [47]: from sklearn.preprocessing import StandardScaler
```

```
In [48]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [51]: prediction=logr.predict(observation)
```

```
print(prediction)
```

```
[28079035]
```

```
In [52]: logr.classes_
```

```
Out[52]: array([28079006, 28079024, 28079035, 28079099], dtype=int64)
```

```
In [53]: logr.score(fs,target_vector)
```

```
Out[53]: 0.8480899292795158
```

```
In [54]: logr.predict_proba(observation)[0][0]
```

```
Out[54]: 2.5638972732451705e-10
```

```
In [55]: logr.predict_proba(observation)
```

```
Out[55]: array([[2.56389727e-10, 3.44199742e-71, 1.00000000e+00, 1.43898646e-13]])
```

Random Forest

```
In [56]: from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: parameters={'max_depth':[1,2,3,4,5],  
                  'min_samples_leaf':[5,10,15,20,25],  
                  'n_estimators':[10,20,30,40,50]  
}
```

```
In [59]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
                      param_grid={'max_depth': [1, 2, 3, 4, 5],  
                                  'min_samples_leaf': [5, 10, 15, 20, 25],  
                                  'n_estimators': [10, 20, 30, 40, 50]},  
                      scoring='accuracy')
```

```
In [60]: grid_search.best_score_
```

```
Out[60]: 0.772919791758581
```

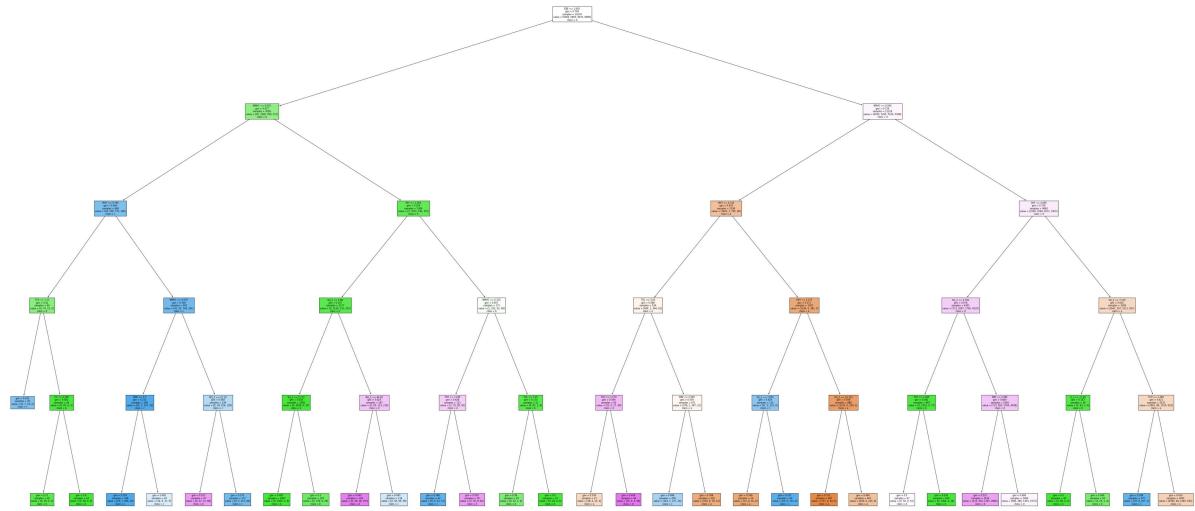
```
In [61]: rfc_best=grid_search.best_estimator_
```

```
In [62]: from sklearn.tree import plot_tree  
  
plt.figure(figsize=(80,40))  
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['a', 'b'])
```

```
Out[62]: [Text(2131.377049180328, 1993.2, 'EBE <= 1.005\ngini = 0.749\nsamples = 14276
\nvalue = [5004, 5605, 5972, 6085]\nclass = d'),
Text(969.639344262295, 1630.8000000000002, 'NMHC <= 0.055\ngini = 0.477\nsamples = 3058\nvalue = [45, 3340, 956, 517]\nclass = b'),
Text(402.4918032786885, 1268.4, 'MXY <= 0.765\ngini = 0.464\nsamples = 669\nvalue = [43, 109, 772, 166]\nclass = c'),
Text(146.36065573770492, 906.0, 'TCH <= 1.23\ngini = 0.42\nsamples = 64\nvalue = [0, 74, 23, 5]\nclass = b'),
Text(73.18032786885246, 543.5999999999999, 'gini = 0.432\nsamples = 20\nvalue = [0, 7, 23, 2]\nclass = c'),
Text(219.54098360655738, 543.5999999999999, 'CO <= 0.285\ngini = 0.082\nsamples = 44\nvalue = [0, 67, 0, 3]\nclass = b'),
Text(146.36065573770492, 181.1999999999982, 'gini = 0.17\nsamples = 20\nvalue = [0, 29, 0, 3]\nclass = b'),
Text(292.72131147540983, 181.1999999999982, 'gini = 0.0\nsamples = 24\nvalue = [0, 38, 0, 0]\nclass = b'),
Text(658.6229508196722, 906.0, 'NMHC <= 0.035\ngini = 0.396\nsamples = 605\nvalue = [43, 35, 749, 161]\nclass = c'),
Text(512.2622950819672, 543.5999999999999, 'BEN <= 1.3\ngini = 0.232\nsamples = 366\nvalue = [43, 1, 523, 33]\nclass = c'),
Text(439.08196721311475, 181.1999999999982, 'gini = 0.203\nsamples = 346\nvalue = [29, 1, 506, 33]\nclass = c'),
Text(585.4426229508197, 181.1999999999982, 'gini = 0.495\nsamples = 20\nvalue = [14, 0, 17, 0]\nclass = c'),
Text(804.983606557377, 543.5999999999999, 'NO_2 <= 21.19\ngini = 0.544\nsamples = 239\nvalue = [0, 34, 226, 128]\nclass = c'),
Text(731.8032786885246, 181.1999999999982, 'gini = 0.527\nsamples = 67\nvalue = [0, 27, 13, 68]\nclass = d'),
Text(878.1639344262295, 181.1999999999982, 'gini = 0.375\nsamples = 172\nvalue = [0, 7, 213, 60]\nclass = c'),
Text(1536.7868852459017, 1268.4, 'PXY <= 1.005\ngini = 0.254\nsamples = 2389
\nvalue = [2, 3231, 184, 351]\nclass = b'),
Text(1244.0655737704917, 906.0, 'SO_2 <= 5.89\ngini = 0.197\nsamples = 2217
\nvalue = [0, 3129, 123, 253]\nclass = b'),
Text(1097.704918032787, 543.5999999999999, 'SO_2 <= 5.115\ngini = 0.024\nsamples = 1950\nvalue = [0, 3034, 0, 37]\nclass = b'),
Text(1024.5245901639344, 181.1999999999982, 'gini = 0.005\nsamples = 1847\nvalue = [0, 2905, 0, 8]\nclass = b'),
Text(1170.8852459016393, 181.1999999999982, 'gini = 0.3\nsamples = 103\nvalue = [0, 129, 0, 29]\nclass = b'),
Text(1390.4262295081967, 543.5999999999999, 'NO_2 <= 44.05\ngini = 0.624\nsamples = 267\nvalue = [0, 95, 123, 216]\nclass = d'),
Text(1317.2459016393443, 181.1999999999982, 'gini = 0.463\nsamples = 149\nvalue = [0, 39, 28, 157]\nclass = d'),
Text(1463.6065573770493, 181.1999999999982, 'gini = 0.645\nsamples = 118\nvalue = [0, 56, 95, 59]\nclass = c'),
Text(1829.5081967213114, 906.0, 'NMHC <= 0.125\ngini = 0.657\nsamples = 172
\nvalue = [2, 102, 61, 98]\nclass = b'),
Text(1683.1475409836066, 543.5999999999999, 'TCH <= 1.245\ngini = 0.628\nsamples = 121\nvalue = [2, 35, 60, 92]\nclass = d'),
Text(1609.967213114754, 181.1999999999982, 'gini = 0.288\nsamples = 42\nvalue = [0, 0, 52, 11]\nclass = c'),
Text(1756.327868852459, 181.1999999999982, 'gini = 0.505\nsamples = 79\nvalue = [2, 35, 8, 81]\nclass = d'),
Text(1975.8688524590164, 543.5999999999999, 'TOL <= 7.25\ngini = 0.173\nsamples = 51\nvalue = [0, 67, 1, 6]\nclass = b'),
Text(1902.688524590164, 181.1999999999982, 'gini = 0.38\nsamples = 20\nvalue = [0, 67, 1, 6]\nclass = b')]
```

```
e = [0, 22, 1, 6]\nclass = b'),  
    Text(2049.0491803278687, 181.19999999999982, 'gini = 0.0\nsamples = 31\nvalue = [0, 45, 0, 0]\nclass = b'),  
    Text(3293.1147540983607, 1630.8000000000002, 'NMHC <= 0.045\ngini = 0.729\nsamples = 11218\nvalue = [4959, 2265, 5016, 5568]\nclass = d'),  
    Text(2707.6721311475408, 1268.4, 'MXY <= 4.145\ngini = 0.476\nsamples = 1538\nvalue = [1603, 1, 745, 96]\nclass = a'),  
    Text(2414.9508196721313, 906.0, 'TOL <= 3.45\ngini = 0.588\nsamples = 536\nvalue = [409, 1, 364, 91]\nclass = a'),  
    Text(2268.590163934426, 543.5999999999999, 'TCH <= 1.175\ngini = 0.589\nsamples = 65\nvalue = [33, 0, 17, 60]\nclass = d'),  
    Text(2195.409836065574, 181.19999999999982, 'gini = 0.539\nsamples = 21\nvalue = [18, 0, 12, 2]\nclass = a'),  
    Text(2341.7704918032787, 181.19999999999982, 'gini = 0.406\nsamples = 44\nvalue = [15, 0, 5, 58]\nclass = d'),  
    Text(2561.311475409836, 543.5999999999999, 'EBE <= 1.545\ngini = 0.539\nsamples = 471\nvalue = [376, 1, 347, 31]\nclass = a'),  
    Text(2488.1311475409834, 181.19999999999982, 'gini = 0.498\nsamples = 269\nvalue = [143, 1, 277, 20]\nclass = c'),  
    Text(2634.4918032786886, 181.19999999999982, 'gini = 0.398\nsamples = 202\nvalue = [233, 0, 70, 11]\nclass = a'),  
    Text(3000.3934426229507, 906.0, 'OXY <= 2.215\ngini = 0.371\nsamples = 1002\nvalue = [1194, 0, 381, 5]\nclass = a'),  
    Text(2854.032786885246, 543.5999999999999, 'SO_2 <= 7.855\ngini = 0.416\nsamples = 114\nvalue = [51, 0, 122, 0]\nclass = c'),  
    Text(2780.8524590163934, 181.19999999999982, 'gini = 0.369\nsamples = 24\nvalue = [31, 0, 10, 0]\nclass = a'),  
    Text(2927.2131147540986, 181.19999999999982, 'gini = 0.257\nsamples = 90\nvalue = [20, 0, 112, 0]\nclass = c'),  
    Text(3146.754098360656, 543.5999999999999, 'SO_2 <= 10.125\ngini = 0.306\nsamples = 888\nvalue = [1143, 0, 259, 5]\nclass = a'),  
    Text(3073.5737704918033, 181.19999999999982, 'gini = 0.121\nsamples = 487\nvalue = [717, 0, 44, 5]\nclass = a'),  
    Text(3219.934426229508, 181.19999999999982, 'gini = 0.446\nsamples = 401\nvalue = [426, 0, 215, 0]\nclass = a'),  
    Text(3878.55737704918, 1268.4, 'OXY <= 4.095\ngini = 0.726\nsamples = 9680\nvalue = [3356, 2264, 4271, 5472]\nclass = d'),  
    Text(3585.8360655737706, 906.0, 'SO_2 <= 5.765\ngini = 0.676\nsamples = 6391\nvalue = [713, 2097, 2760, 4535]\nclass = d'),  
    Text(3439.4754098360654, 543.5999999999999, 'TCH <= 1.265\ngini = 0.092\nsamples = 997\nvalue = [0, 1514, 0, 77]\nclass = b'),  
    Text(3366.2950819672133, 181.19999999999982, 'gini = 0.5\nsamples = 67\nvalue = [0, 50, 0, 51]\nclass = d'),  
    Text(3512.655737704918, 181.19999999999982, 'gini = 0.034\nsamples = 930\nvalue = [0, 1464, 0, 26]\nclass = b'),  
    Text(3732.1967213114754, 543.5999999999999, 'EBE <= 2.285\ngini = 0.609\nsamples = 5394\nvalue = [713, 583, 2760, 4458]\nclass = d'),  
    Text(3659.0163934426228, 181.19999999999982, 'gini = 0.527\nsamples = 2914\nvalue = [133, 303, 1307, 2885]\nclass = d'),  
    Text(3805.377049180328, 181.19999999999982, 'gini = 0.669\nsamples = 2480\nvalue = [580, 280, 1453, 1573]\nclass = d'),  
    Text(4171.2786885245905, 906.0, 'SO_2 <= 7.125\ngini = 0.632\nsamples = 3289\nvalue = [2643, 167, 1511, 937]\nclass = a'),  
    Text(4024.9180327868853, 543.5999999999999, 'O_3 <= 11.99\ngini = 0.147\nsamples = 56\nvalue = [2, 83, 1, 4]\nclass = b'),  
    Text(3951.7377049180327, 181.19999999999982, 'gini = 0.0\nsamples = 36\nvalue = [0, 60, 0, 0]\nclass = b'),
```

```
Text(4098.098360655737, 181.19999999999982, 'gini = 0.389\nsamples = 20\nvalue = [2, 23, 1, 4]\nclass = b'),
Text(4317.639344262295, 543.5999999999999, 'TCH <= 1.285\n gini = 0.621\nsamples = 3233\nvalue = [2641, 84, 1510, 933]\nclass = a'),
Text(4244.459016393443, 181.19999999999982, 'gini = 0.289\nsamples = 172\nvalue = [47, 0, 227, 1]\nclass = c'),
Text(4390.819672131148, 181.19999999999982, 'gini = 0.614\nsamples = 3061\nvalue = [2594, 84, 1283, 932]\nclass = a')]
```



Conclusion

Accuracy

Linear Regression:0.20410773820018646

Ridge Regression:0.20390330994003047

Lasso Regression:0.061940553181437985

ElasticNet Regression:0.08795788332481069

Logistic Regression:0.8480899292795158

Random Forest:0.772919791758581

Logistic Regression is suitable for this dataset

