

Problem Statement

A real estate agent want help to predict the house price for regions in USA.He gave us the dataset to work on to use linear regression model.Create a model that helps him to estimate of what the house would sell for

Import libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: # To import dataset
df=pd.read_csv('20_states.csv')
df
```

Out[2]:

	id	name	country_id	country_code	country_name	state_code	type	latitude	
	0	3901	Badakhshan	1	AF	Afghanistan	BDS	NaN	36.734772
	1	3871	Badghis	1	AF	Afghanistan	BDG	NaN	35.167134
	2	3875	Baghlan	1	AF	Afghanistan	BGL	NaN	36.178903
	3	3884	Balkh	1	AF	Afghanistan	BAL	NaN	36.755060
	4	3872	Bamyan	1	AF	Afghanistan	BAM	NaN	34.810007

5072	1953	Mashonaland West Province	247	ZW	Zimbabwe	MW	NaN	-17.485103	
5073	1960	Masvingo Province	247	ZW	Zimbabwe	MV	NaN	-20.624151	
5074	1954	Matabeleland North Province	247	ZW	Zimbabwe	MN	NaN	-18.533157	
5075	1952	Matabeleland South Province	247	ZW	Zimbabwe	MS	NaN	-21.052337	
5076	1957	Midlands Province	247	ZW	Zimbabwe	MI	NaN	-19.055201	

5077 rows × 9 columns



```
In [3]: # To display top 10 rows
df.head(10)
```

Out[3]:

	id	name	country_id	country_code	country_name	state_code	type	latitude	longitude
0	3901	Badakhshan	1	AF	Afghanistan	BDS	NaN	36.734772	70.8
1	3871	Badghis	1	AF	Afghanistan	BDG	NaN	35.167134	63.7
2	3875	Baghlan	1	AF	Afghanistan	BGL	NaN	36.178903	68.7
3	3884	Balkh	1	AF	Afghanistan	BAL	NaN	36.755060	66.8
4	3872	Bamyan	1	AF	Afghanistan	BAM	NaN	34.810007	67.8
5	3892	Daykundi	1	AF	Afghanistan	DAY	NaN	33.669495	66.0
6	3899	Farah	1	AF	Afghanistan	FRA	NaN	32.495328	62.2
7	3889	Faryab	1	AF	Afghanistan	FYB	NaN	36.079561	64.9
8	3870	Ghazni	1	AF	Afghanistan	GHA	NaN	33.545059	68.4
9	3888	Ghōr	1	AF	Afghanistan	GHO	NaN	34.099578	64.9

Data Cleaning and Pre-Processing

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5077 entries, 0 to 5076
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id               5077 non-null   int64
1   name             5077 non-null   object
2   country_id       5077 non-null   int64
3   country_code     5063 non-null   object
4   country_name     5077 non-null   object
5   state_code       5072 non-null   object
6   type             1597 non-null   object
7   latitude         5008 non-null   float64
8   longitude        5008 non-null   float64
dtypes: float64(2), int64(2), object(5)
memory usage: 357.1+ KB
```

```
In [5]: df.describe()
```

```
Out[5]:
```

	id	country_id	latitude	longitude
count	5077.000000	5077.000000	5008.000000	5008.000000
mean	2609.765413	133.467599	27.576415	17.178713
std	1503.376799	72.341160	22.208161	61.269334
min	1.000000	1.000000	-54.805400	-178.116500
25%	1324.000000	74.000000	11.399747	-3.943859
50%	2617.000000	132.000000	34.226432	17.501792
75%	3905.000000	201.000000	45.802822	41.919647
max	5220.000000	248.000000	77.874972	179.852222

```
In [6]: df.columns
```

```
Out[6]: Index(['id', 'name', 'country_id', 'country_code', 'country_name',  
              'state_code', 'type', 'latitude', 'longitude'],  
              dtype='object')
```


```
In [7]: a = df.dropna(axis='columns')  
a.columns
```

```
Out[7]: Index(['id', 'name', 'country_id', 'country_name'], dtype='object')
```

EDA and Visualization

```
In [8]: sns.pairplot(a)
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x22239f86af0>
```




```
In [9]: sns.distplot(a['country_id'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```


```
Out[9]: <AxesSubplot:xlabel='country_id', ylabel='Density'>
```



```
In [10]: a1=a[['id', 'country_id']]
```

```
In [11]: sns.heatmap(a1.corr())
```

```
Out[11]: <AxesSubplot:>
```



To Train the Model - Model Building

We are going to train Linear Regression model; We need to split out data into two variables x and y where x is independent variable (input) and y is dependent on x (output). We could ignore address column as it is not required for our model.

```
In [12]: x=a1[['id']]  
y=a1['country_id']
```

To split my dataset into training and test data

```
In [13]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [14]: from sklearn.linear_model import LinearRegression  
  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[14]: LinearRegression()
```

```
In [15]: print(lr.intercept_)
```

```
127.57310018216988
```


```
In [16]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[16]:

Co-efficient	
id	0.002029

```
In [17]: prediction=lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[17]: <matplotlib.collections.PathCollection at 0x2223b38a820>



```
In [18]: print(lr.score(x_test,y_test))
```

0.0017083297723273771

```
In [19]: from sklearn.linear_model import Ridge,Lasso
```

```
In [20]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[20]: Ridge(alpha=10)

```
In [21]: rr.score(x_train,y_train)
```

Out[21]: 0.0017342131031994334

```
In [22]: rr.score(x_test,y_test)
```

Out[22]: 0.0017083297715219103

```
In [23]: rr.score(x_test,y_test)
```

Out[23]: 0.0017083297715219103

```
In [24]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[24]: Lasso(alpha=10)

```
In [25]: la.score(x_test,y_test)
```

Out[25]: 0.001706910407451523

```
In [26]: from sklearn.linear_model import ElasticNet
en = ElasticNet()
en.fit(x_train,y_train)
```

Out[26]: ElasticNet()

```
In [27]: print(en.coef_)
```

[0.0020287]

```
In [28]: print(en.intercept_)
```

127.5736789283807

```
In [29]: print(en.predict(x_test))
```

[137.43316486 137.48996848 130.72019388 ... 127.95507468 136.5466226
134.85265743]

```
In [30]: print(en.score(x_test,y_test))
```

0.0017082590891673854

Evaluation Metrics

```
In [31]: from sklearn import metrics
print("Mean Absolytre Error:",metrics.mean_absolute_error(y_test,prediction))
print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,pre
```

Mean Absolytre Error: 59.86992832773104
Mean Squared Error: 4933.335203113695
Root Mean Squared Error: 70.23770499606101

```
In [32]: import pickle
```

```
In [33]: filename='prediction2'
pickle.dump(lr,open(filename,'wb'))
```

In []: