

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
In [2]: df=pd.read_csv("ionosphere.csv")
df
```

Out[2]:

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.03760	...	-0.51171	0.41078	-0.46168
0	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	...	-0.26569	-0.20468	-0.18401
1	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	...	-0.40220	0.58984	-0.22145
2	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	...	0.90695	0.51613	1.00000
3	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	...	-0.65158	0.13290	-0.53206
4	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	...	-0.01535	-0.03240	0.09223
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
345	1	0	0.83508	0.08298	0.73739	-0.14706	0.84349	-0.05567	0.90441	-0.04622	...	-0.04202	0.83479	0.00145
346	1	0	0.95113	0.00419	0.95183	-0.02723	0.93438	-0.01920	0.94590	0.01606	...	0.01361	0.93522	0.04923
347	1	0	0.94701	-0.00034	0.93207	-0.03227	0.95177	-0.03431	0.95584	0.02446	...	0.03193	0.92489	0.02506
348	1	0	0.90608	-0.01657	0.98122	-0.01989	0.95691	-0.03646	0.85746	0.00110	...	-0.02099	0.89147	-0.07714
349	1	0	0.84710	0.13533	0.73638	-0.06151	0.87873	0.08260	0.88928	-0.09139	...	-0.15114	0.81147	-0.04823

350 rows × 35 columns



```
In [3]: df.head()
```

Out[3]:

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.03760	...	-0.51171	0.41078	-0.46168
0	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	...	-0.26569	-0.20468	-0.18401
1	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	...	-0.40220	0.58984	-0.22145
2	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	...	0.90695	0.51613	1.00000
3	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	...	-0.65158	0.13290	-0.53206
4	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	...	-0.01535	-0.03240	0.09223

5 rows × 35 columns



## Data Cleaning and Data Preprocessing

In [4]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 350 entries, 0 to 349
Data columns (total 35 columns):
#   Column      Non-Null Count  Dtype
---  -
0    1           350 non-null    int64
1    0           350 non-null    int64
2    0.99539     350 non-null    float64
3    -0.05889    350 non-null    float64
4    0.85243     350 non-null    float64
5    0.02306     350 non-null    float64
6    0.83398     350 non-null    float64
7    -0.37708    350 non-null    float64
8    1.1         350 non-null    float64
9    0.03760     350 non-null    float64
10   0.85243.1   350 non-null    float64
11   -0.17755    350 non-null    float64
12   0.59755     350 non-null    float64
13   -0.44945    350 non-null    float64
14   0.60536     350 non-null    float64
15   -0.38223    350 non-null    float64
16   0.84356     350 non-null    float64
17   -0.38542    350 non-null    float64
18   0.58212     350 non-null    float64
19   -0.32192    350 non-null    float64
20   0.56971     350 non-null    float64
21   -0.29674    350 non-null    float64
22   0.36946     350 non-null    float64
23   -0.47357    350 non-null    float64
24   0.56811     350 non-null    float64
25   -0.51171    350 non-null    float64
26   0.41078     350 non-null    float64
27   -0.46168    350 non-null    float64
28   0.21266     350 non-null    float64
29   -0.34090    350 non-null    float64
30   0.42267     350 non-null    float64
31   -0.54487    350 non-null    float64
32   0.18641     350 non-null    float64
33   -0.45300    350 non-null    float64
34   g           350 non-null    object
dtypes: float64(32), int64(2), object(1)
memory usage: 95.8+ KB


```

In [5]: `df.describe()`

Out[5]:

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.
<b>count</b>	350.000000	350.0	350.000000	350.000000	350.000000	350.000000	350.000000	350.000000	350.000000	350.0
<b>mean</b>	0.891429	0.0	0.640330	0.044667	0.600350	0.116154	0.549284	0.120779	0.510453	0.1
<b>std</b>	0.311546	0.0	0.498059	0.442032	0.520431	0.461443	0.493124	0.520816	0.507117	0.4
<b>min</b>	0.000000	0.0	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	-1.0
<b>25%</b>	1.000000	0.0	0.471517	-0.065388	0.412555	-0.024868	0.209105	-0.053483	0.086785	-0.0
<b>50%</b>	1.000000	0.0	0.870795	0.016700	0.808620	0.021170	0.728000	0.015085	0.682430	0.0
<b>75%</b>	1.000000	0.0	1.000000	0.194727	1.000000	0.335317	0.970445	0.451572	0.950555	0.5
<b>max</b>	1.000000	0.0	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.0

8 rows × 34 columns



In [6]: `df.columns`

Out[6]: Index(['1', '0', '0.99539', '-0.05889', '0.85243', '0.02306', '0.83398',  
'-0.37708', '1.1', '0.03760', '0.85243.1', '-0.17755', '0.59755',  
'-0.44945', '0.60536', '-0.38223', '0.84356', '-0.38542', '0.58212',  
'-0.32192', '0.56971', '-0.29674', '0.36946', '-0.47357', '0.56811',  
'-0.51171', '0.41078', '-0.46168', '0.21266', '-0.34090', '0.42267',  
'-0.54487', '0.18641', '-0.45300', 'g'],  
dtype='object')

In [7]: `feature_matrix = df.iloc[:,0:34]`  
`target_vector = df.iloc[:,-1]`

In [8]: `fs = StandardScaler().fit_transform(feature_matrix)`  
`logr = LogisticRegression()`  
`logr.fit(fs,target_vector)`

Out[8]: LogisticRegression()

```
In [9]: observation=[[1.0,0.0,1.0,-0.18829,0.93035,  
-0.36156,  
-0.10868,  
-0.93597,  
1.0,  
-0.04549,  
0.50874,  
-0.67743,  
0.34432,  
-0.69707,  
-0.51685,  
-0.97515,  
0.05499,  
-0.62237,  
0.33109,  
-1.0,  
-0.13151,  
-0.453,  
-0.18056,  
-0.35734,  
-0.20332,  
-0.26569,  
-0.20468,  
-0.18401,  
-0.1904,  
-0.11593,  
-0.16626,  
-0.06288,  
-0.13738,  
-0.02447]]  
prediction = logr.predict(observation)  
print(prediction)
```

```
['g']
```

```
In [10]: logr.classes_
```

```
Out[10]: array(['b', 'g'], dtype=object)
```

```
In [11]: logr.predict_proba(observation)
```

```
Out[11]: array([[0.07006552, 0.92993448]])
```

## Random Forest

```
In [12]: df['g'].value_counts()
```

```
Out[12]: g      224  
b       126  
Name: g, dtype: int64
```

```
In [13]: x=df.drop('g', axis=1)  
y=df['g']
```

```
In [14]: g1={"g":{"g":1, "b":2}}
df=df.replace(g1)
df
```

```
Out[14]:
```

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.03760	...	-0.51171	0.41078	-0.4611
0	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	...	-0.26569	-0.20468	-0.1841
1	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	...	-0.40220	0.58984	-0.2211
2	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	...	0.90695	0.51613	1.0000
3	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	...	-0.65158	0.13290	-0.5321
4	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	...	-0.01535	-0.03240	0.0921
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
345	1	0	0.83508	0.08298	0.73739	-0.14706	0.84349	-0.05567	0.90441	-0.04622	...	-0.04202	0.83479	0.0011
346	1	0	0.95113	0.00419	0.95183	-0.02723	0.93438	-0.01920	0.94590	0.01606	...	0.01361	0.93522	0.0491
347	1	0	0.94701	-0.00034	0.93207	-0.03227	0.95177	-0.03431	0.95584	0.02446	...	0.03193	0.92489	0.0251
348	1	0	0.90608	-0.01657	0.98122	-0.01989	0.95691	-0.03646	0.85746	0.00110	...	-0.02099	0.89147	-0.0771
349	1	0	0.84710	0.13533	0.73638	-0.06151	0.87873	0.08260	0.88928	-0.09139	...	-0.15114	0.81147	-0.0481

350 rows × 35 columns



```
In [15]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
In [16]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[16]: RandomForestClassifier()
```

```
In [21]: parameters = {'max_depth':[1,2,3,4,5], 'min_samples_leaf':[5,10,15,20,25],
                        'n_estimators': [10,20,30,40,50]}
}
```

```
In [22]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[22]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 3, 4, 5],
                                   'min_samples_leaf': [5, 10, 15, 20, 25],
                                   'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

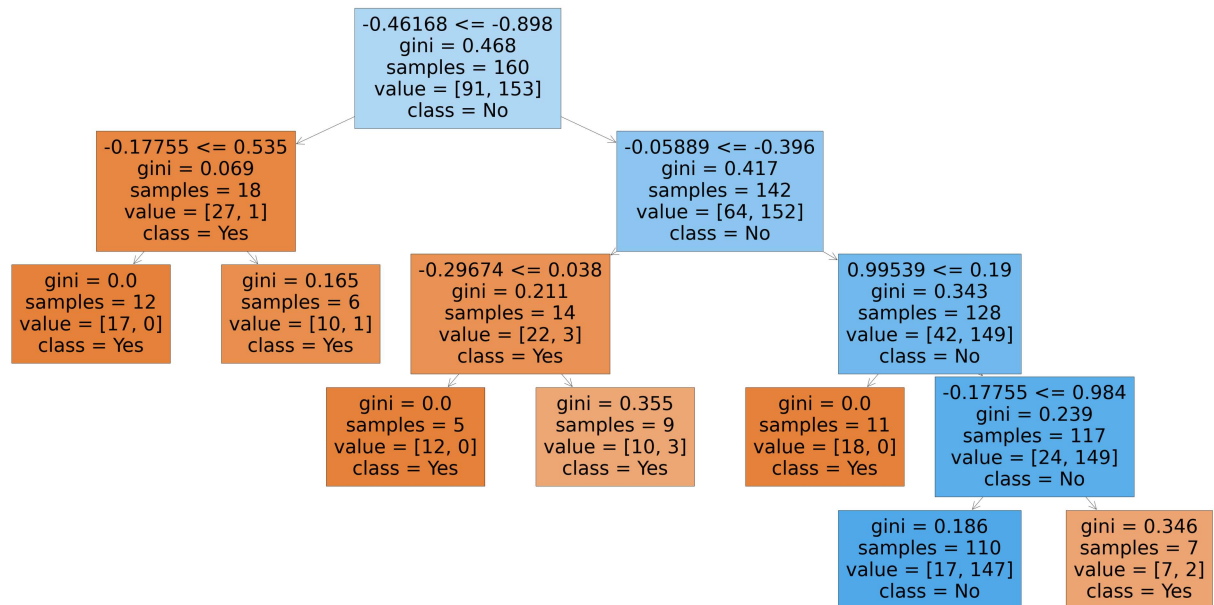
```
In [23]: grid_search.best_score_
```

```
Out[23]: 0.930327868852459
```

```
In [24]: rfc_best = grid_search.best_estimator_
```

```
In [27]: from sklearn.tree import plot_tree
plt.figure(figsize = (80,40,))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'],filled=True)
```

```
Out[27]: [Text(1674.0, 1956.96, '-0.46168 <= -0.898\ngini = 0.468\nsamples = 160\nvalue = [91, 153]\nnclass = No'),
Text(744.0, 1522.0800000000002, '-0.17755 <= 0.535\ngini = 0.069\nsamples = 18\nvalue = [27, 1]\nnclass = Yes'),
Text(372.0, 1087.2, 'gini = 0.0\nsamples = 12\nvalue = [17, 0]\nnclass = Yes'),
Text(1116.0, 1087.2, 'gini = 0.165\nsamples = 6\nvalue = [10, 1]\nnclass = Yes'),
Text(2604.0, 1522.0800000000002, '-0.05889 <= -0.396\ngini = 0.417\nsamples = 142\nvalue = [64, 152]\nnclass = No'),
Text(1860.0, 1087.2, '-0.29674 <= 0.038\ngini = 0.211\nsamples = 14\nvalue = [22, 3]\nnclass = Yes'),
Text(1488.0, 652.3200000000002, 'gini = 0.0\nsamples = 5\nvalue = [12, 0]\nnclass = Yes'),
Text(2232.0, 652.3200000000002, 'gini = 0.355\nsamples = 9\nvalue = [10, 3]\nnclass = Yes'),
Text(3348.0, 1087.2, '0.99539 <= 0.19\ngini = 0.343\nsamples = 128\nvalue = [42, 149]\nnclass = No'),
Text(2976.0, 652.3200000000002, 'gini = 0.0\nsamples = 11\nvalue = [18, 0]\nnclass = Yes'),
Text(3720.0, 652.3200000000002, '-0.17755 <= 0.984\ngini = 0.239\nsamples = 117\nvalue = [24, 149]\nnclass = No'),
Text(3348.0, 217.44000000000005, 'gini = 0.186\nsamples = 110\nvalue = [17, 147]\nnclass = No'),
Text(4092.0, 217.44000000000005, 'gini = 0.346\nsamples = 7\nvalue = [7, 2]\nnclass = Yes')]
```



In [ ]: