

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: from sklearn.linear_model import LogisticRegression
```

```
In [3]: df=pd.read_csv("C5 health.csv").dropna()

df
```

```
Out[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148	72	35	0	33.6	0.62
1	1	85	66	29	0	26.6	0.35
2	8	183	64	0	0	23.3	0.67
3	1	89	66	23	94	28.1	0.16
4	0	137	40	35	168	43.1	2.28
...	...	...	...	...	...	...	...
763	10	101	76	48	180	32.9	0.17
764	2	122	70	27	0	36.8	0.34
765	5	121	72	23	112	26.2	0.24
766	1	126	60	0	0	30.1	0.34
767	1	93	70	31	0	30.4	0.31

768 rows × 9 columns



```
In [4]: df.dropna(inplace=True)
```

In [5]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Pregnancies           768 non-null    int64
 1   Glucose               768 non-null    int64
 2   BloodPressure         768 non-null    int64
 3   SkinThickness         768 non-null    int64
 4   Insulin               768 non-null    int64
 5   BMI                  768 non-null    float64
 6   DiabetesPedigreeFunction 768 non-null    float64
 7   Age                  768 non-null    int64
 8   Outcome               768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 60.0 KB
```

In [6]: feature\_matrix = df[['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'DiabetesPedigreeFunction', 'Age', 'Outcome']]  
target\_vector = df['Outcome']

In [7]: feature\_matrix.shape

Out[7]: (768, 8)

In [8]: target\_vector.shape

Out[8]: (768,)

In [9]: from sklearn.preprocessing import StandardScaler

In [10]: fs = StandardScaler().fit\_transform(feature\_matrix)

In [11]: logr = LogisticRegression()  
logr.fit(fs, target\_vector)

Out[11]: LogisticRegression()

In [12]: feature\_matrix.shape

Out[12]: (768, 8)

In [13]: target\_vector.shape

Out[13]: (768,)

In [14]: from sklearn.preprocessing import StandardScaler

In [15]: fs = StandardScaler().fit\_transform(feature\_matrix)



```
In [23]: g1={'Outcome':{'1':1, "0":2}}
df=df.replace(g1)
df
```

```
Out[23]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFuncio
0	6	148	72	35	0	33.6	0.62
1	1	85	66	29	0	26.6	0.35
2	8	183	64	0	0	23.3	0.67
3	1	89	66	23	94	28.1	0.16
4	0	137	40	35	168	43.1	2.28
...	...	...	...	...	...	...	...
763	10	101	76	48	180	32.9	0.17
764	2	122	70	27	0	36.8	0.34
765	5	121	72	23	112	26.2	0.24
766	1	126	60	0	0	30.1	0.34
767	1	93	70	31	0	30.4	0.31

768 rows × 9 columns

```
In [24]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
In [25]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[25]: RandomForestClassifier()
```

```
In [26]: parameters = {'max_depth':[1,2,3,4,5], 'min_samples_leaf':[5,10,15,20,25],
                        'n_estimators': [10,20,30,40,50]}
}
```

```
In [27]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[27]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 3, 4, 5],
                                   'min_samples_leaf': [5, 10, 15, 20, 25],
                                   'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

```
In [28]: grid_search.best_score_
```

```
Out[28]: 0.7821256172668257
```

```
In [29]: rfc_best = grid_search.best_estimator_
```

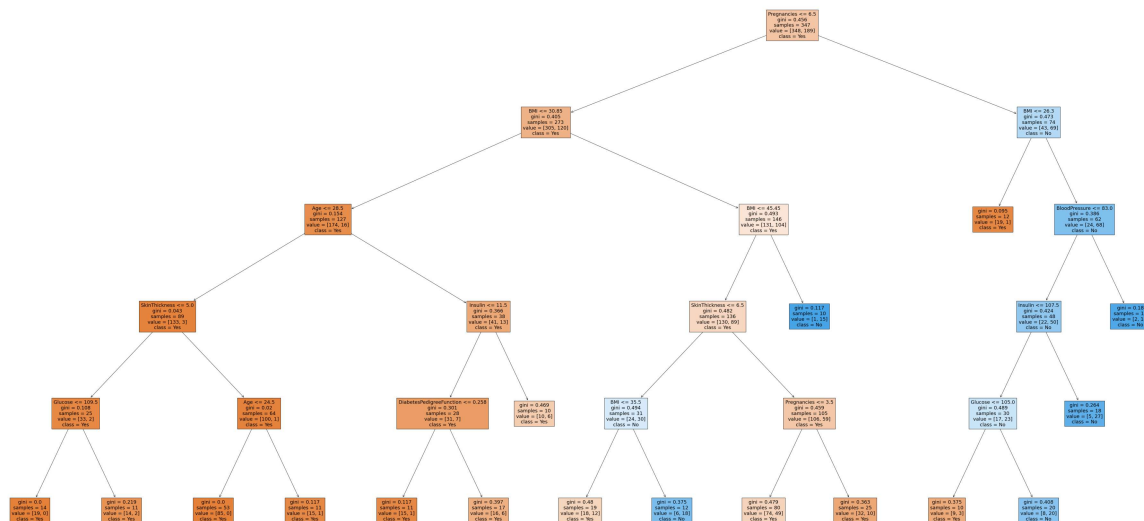
```
In [30]: from sklearn.tree import plot_tree  
plt.figure(figsize = (80,40,))  
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'])
```

```

Out[30]: [Text(3026.076923076923, 1993.2, 'Pregnancies <= 6.5\ngini = 0.456\nsamples =
347\nvalue = [348, 189]\nclass = Yes'),
Text(2103.230769230769, 1630.8000000000002, 'BMI <= 30.85\ngini = 0.405\nsam
ples = 273\nvalue = [305, 120]\nclass = Yes'),
Text(1287.6923076923076, 1268.4, 'Age <= 28.5\ngini = 0.154\nsamples = 127\n
value = [174, 16]\nclass = Yes'),
Text(686.7692307692307, 906.0, 'SkinThickness <= 5.0\ngini = 0.043\nsamples
= 89\nvalue = [133, 3]\nclass = Yes'),
Text(343.38461538461536, 543.5999999999999, 'Glucose <= 109.5\ngini = 0.108
\nsamples = 25\nvalue = [33, 2]\nclass = Yes'),
Text(171.69230769230768, 181.19999999999982, 'gini = 0.0\nsamples = 14\nvalu
e = [19, 0]\nclass = Yes'),
Text(515.0769230769231, 181.19999999999982, 'gini = 0.219\nsamples = 11\nval
ue = [14, 2]\nclass = Yes'),
Text(1030.1538461538462, 543.5999999999999, 'Age <= 24.5\ngini = 0.02\nsampl
es = 64\nvalue = [100, 1]\nclass = Yes'),
Text(858.4615384615383, 181.19999999999982, 'gini = 0.0\nsamples = 53\nvalue
= [85, 0]\nclass = Yes'),
Text(1201.8461538461538, 181.19999999999982, 'gini = 0.117\nsamples = 11\nva
lue = [15, 1]\nclass = Yes'),
Text(1888.6153846153845, 906.0, 'Insulin <= 11.5\ngini = 0.366\nsamples = 38
\nvalue = [41, 13]\nclass = Yes'),
Text(1716.9230769230767, 543.5999999999999, 'DiabetesPedigreeFunction <= 0.2
58\ngini = 0.301\nsamples = 28\nvalue = [31, 7]\nclass = Yes'),
Text(1545.230769230769, 181.19999999999982, 'gini = 0.117\nsamples = 11\nval
ue = [15, 1]\nclass = Yes'),
Text(1888.6153846153845, 181.19999999999982, 'gini = 0.397\nsamples = 17\nva
lue = [16, 6]\nclass = Yes'),
Text(2060.3076923076924, 543.5999999999999, 'gini = 0.469\nsamples = 10\nval
ue = [10, 6]\nclass = Yes'),
Text(2918.7692307692305, 1268.4, 'BMI <= 45.45\ngini = 0.493\nsamples = 146
\nvalue = [131, 104]\nclass = Yes'),
Text(2747.076923076923, 906.0, 'SkinThickness <= 6.5\ngini = 0.482\nsamples
= 136\nvalue = [130, 89]\nclass = Yes'),
Text(2403.6923076923076, 543.5999999999999, 'BMI <= 35.5\ngini = 0.494\nsamp
les = 31\nvalue = [24, 30]\nclass = No'),
Text(2232.0, 181.19999999999982, 'gini = 0.48\nsamples = 19\nvalue = [18, 1
2]\nclass = Yes'),
Text(2575.3846153846152, 181.19999999999982, 'gini = 0.375\nsamples = 12\nva
lue = [6, 18]\nclass = No'),
Text(3090.461538461538, 543.5999999999999, 'Pregnancies <= 3.5\ngini = 0.459
\nsamples = 105\nvalue = [106, 59]\nclass = Yes'),
Text(2918.7692307692305, 181.19999999999982, 'gini = 0.479\nsamples = 80\nva
lue = [74, 49]\nclass = Yes'),
Text(3262.1538461538457, 181.19999999999982, 'gini = 0.363\nsamples = 25\nva
lue = [32, 10]\nclass = Yes'),
Text(3090.461538461538, 906.0, 'gini = 0.117\nsamples = 10\nvalue = [1, 15]
\nclass = No'),
Text(3948.9230769230767, 1630.8000000000002, 'BMI <= 26.3\ngini = 0.473\nsam
ples = 74\nvalue = [43, 69]\nclass = No'),
Text(3777.230769230769, 1268.4, 'gini = 0.095\nsamples = 12\nvalue = [19, 1]
\nclass = Yes'),
Text(4120.615384615385, 1268.4, 'BloodPressure <= 83.0\ngini = 0.386\nsampl
es = 62\nvalue = [24, 68]\nclass = No'),
Text(3948.9230769230767, 906.0, 'Insulin <= 107.5\ngini = 0.424\nsamples = 4
8\nvalue = [22, 50]\nclass = No'),
Text(3777.230769230769, 543.5999999999999, 'Glucose <= 105.0\ngini = 0.489\n

```

```
samples = 30\nvalue = [17, 23]\nclass = No'),
Text(3605.5384615384614, 181.19999999999982, 'gini = 0.375\nsamples = 10\nvalue = [9, 3]\nclass = Yes'),
Text(3948.9230769230767, 181.19999999999982, 'gini = 0.408\nsamples = 20\nvalue = [8, 20]\nclass = No'),
Text(4120.615384615385, 543.5999999999999, 'gini = 0.264\nsamples = 18\nvalue = [5, 27]\nclass = No'),
Text(4292.307692307692, 906.0, 'gini = 0.18\nsamples = 14\nvalue = [2, 18]\nclass = No')]
```



Loading [MathJax]/extensions/Safe.js