

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: from sklearn.linear_model import LogisticRegression
```

```
In [3]: df=pd.read_csv("C6 bmi csv").dropna()
df
```

```
Out[3]:
```

	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3
...
495	Female	150	153	5
496	Female	184	121	4
497	Female	141	136	5
498	Male	150	95	5
499	Male	173	131	5

500 rows × 4 columns

```
In [4]: df.dropna(inplace=True)
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 500 entries, 0 to 499
Data columns (total 4 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Gender  500 non-null      object
1   Height  500 non-null      int64
2   Weight  500 non-null      int64
3   Index   500 non-null      int64
dtypes: int64(3), object(1)
memory usage: 19.5+ KB
```

```
In [6]: feature_matrix = df[['Index', 'Height', 'Weight']]
target_vector = df['Gender']
```

```
In [7]: feature_matrix.shape
```

```
Out[7]: (500, 3)
```

```
In [8]: target_vector.shape
```

```
Out[8]: (500,)
```

```
In [9]: from sklearn.preprocessing import StandardScaler
```

```
In [10]: fs = StandardScaler().fit_transform(feature_matrix)
```

```
In [11]: logr = LogisticRegression()  
logr.fit(fs,target_vector)
```

```
Out[11]: LogisticRegression()
```

```
In [12]: feature_matrix.shape
```

```
Out[12]: (500, 3)
```

```
In [13]: target_vector.shape
```

```
Out[13]: (500,)
```

```
In [14]: from sklearn.preprocessing import StandardScaler
```

```
In [15]: fs = StandardScaler().fit_transform(feature_matrix)
```

```
In [16]: logr = LogisticRegression()  
logr.fit(fs,target_vector)
```

```
Out[16]: LogisticRegression()
```

```
In [17]: observation=df[['Index','Height','Weight']]
```



```
In [23]: g1={'Gender':{'Male':1, "Female":2}}
df=df.replace(g1)
df
```

```
Out[23]:
```

	Gender	Height	Weight	Index
0	1	174	96	4
1	1	189	87	2
2	2	185	110	4
3	2	195	104	3
4	1	149	61	3
...
495	2	150	153	5
496	2	184	121	4
497	2	141	136	5
498	1	150	95	5
499	1	173	131	5

500 rows × 4 columns

```
In [24]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
In [25]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[25]: RandomForestClassifier()
```

```
In [26]: parameters = {'max_depth':[1,2,3,4,5], 'min_samples_leaf':[5,10,15,20,25],
                        'n_estimators': [10,20,30,40,50]}
}
```

```
In [27]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="a
grid_search.fit(x_train,y_train)
```

```
Out[27]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 3, 4, 5],
                                   'min_samples_leaf': [5, 10, 15, 20, 25],
                                   'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

```
In [28]: grid_search.best_score_
```

```
Out[28]: 0.5
```

```
In [29]: rfc_best = grid_search.best_estimator_
```

```
In [30]: from sklearn.tree import plot_tree  
plt.figure(figsize = (80,40,))  
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['Yes','No'])
```

```
Out[30]: [Text(2232.0, 1630.8000000000002, 'Index <= 1.5\ngini = 0.5\nsamples = 220\nvalue = [180, 170]\n\nclass = Yes'),  
Text(1116.0, 543.5999999999999, 'gini = 0.278\nsamples = 9\nvalue = [3, 15]\n\nclass = No'),  
Text(3348.0, 543.5999999999999, 'gini = 0.498\nsamples = 211\nvalue = [177, 155]\n\nclass = Yes')]
```

