

Problem Statement

A real estate agent want help to predict the house price for regions in USA.He gave us the dataset to work on to use linear regression model.Create a model that helps him to estimate of what the house would sell for

Import libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: # To import dataset
df=pd.read_csv('drug csv')
df
```

```
Out[2]:
```

| | Age | Sex | BP | Cholesterol | Na_to_K | Drug |
|-----|-----|-----|--------|-------------|---------|-------|
| 0 | 23 | F | HIGH | HIGH | 25.355 | drugY |
| 1 | 47 | M | LOW | HIGH | 13.093 | drugC |
| 2 | 47 | M | LOW | HIGH | 10.114 | drugC |
| 3 | 28 | F | NORMAL | HIGH | 7.798 | drugX |
| 4 | 61 | F | LOW | HIGH | 18.043 | drugY |
| ... | ... | ... | ... | ... | ... | ... |
| 195 | 56 | F | LOW | HIGH | 11.567 | drugC |
| 196 | 16 | M | LOW | HIGH | 12.006 | drugC |
| 197 | 52 | M | NORMAL | HIGH | 9.894 | drugX |
| 198 | 23 | M | NORMAL | NORMAL | 14.020 | drugX |
| 199 | 40 | F | LOW | NORMAL | 11.349 | drugX |

200 rows × 6 columns

```
In [3]: # To display top 10 rows
df.head(10)
```

```
Out[3]:
```

| | Age | Sex | BP | Cholesterol | Na_to_K | Drug |
|---|-----|-----|--------|-------------|---------|-------|
| 0 | 23 | F | HIGH | HIGH | 25.355 | drugY |
| 1 | 47 | M | LOW | HIGH | 13.093 | drugC |
| 2 | 47 | M | LOW | HIGH | 10.114 | drugC |
| 3 | 28 | F | NORMAL | HIGH | 7.798 | drugX |
| 4 | 61 | F | LOW | HIGH | 18.043 | drugY |
| 5 | 22 | F | NORMAL | HIGH | 8.607 | drugX |
| 6 | 49 | F | NORMAL | HIGH | 16.275 | drugY |
| 7 | 41 | M | LOW | HIGH | 11.037 | drugC |
| 8 | 60 | M | NORMAL | HIGH | 15.171 | drugY |
| 9 | 43 | M | LOW | NORMAL | 19.368 | drugY |

Data Cleaning and Pre-Processing

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age             200 non-null   int64
1   Sex             200 non-null   object
2   BP              200 non-null   object
3   Cholesterol     200 non-null   object
4   Na_to_K         200 non-null   float64
5   Drug            200 non-null   object
dtypes: float64(1), int64(1), object(4)
memory usage: 9.5+ KB
```

```
In [5]: df.describe()
```

```
Out[5]:
```

| | Age | Na_to_K |
|-------|------------|------------|
| count | 200.000000 | 200.000000 |
| mean | 44.315000 | 16.084485 |
| std | 16.544315 | 7.223956 |
| min | 15.000000 | 6.269000 |
| 25% | 31.000000 | 10.445500 |
| 50% | 45.000000 | 13.936500 |
| 75% | 58.000000 | 19.380000 |
| max | 74.000000 | 38.247000 |

```
In [6]: df.columns
```

```
Out[6]: Index(['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K', 'Drug'], dtype='object')
```


```
In [7]: a = df.dropna(axis='columns')  
a.columns
```

```
Out[7]: Index(['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K', 'Drug'], dtype='object')
```

EDA and Visualization

```
In [8]: sns.pairplot(a)
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x18efd21a910>
```

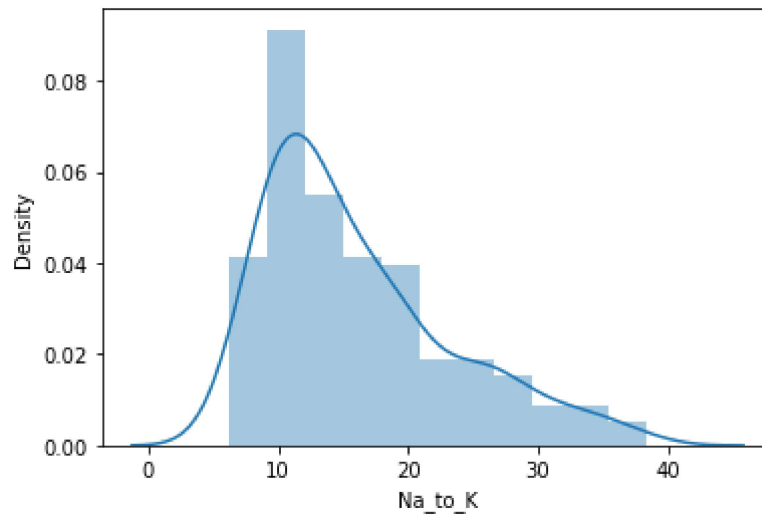


```
In [9]: sns.distplot(a['Na_to_K'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

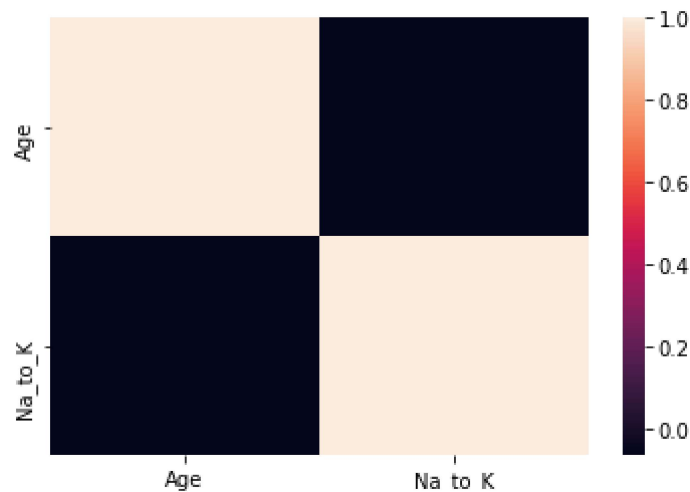
```
Out[9]: <AxesSubplot:xlabel='Na_to_K', ylabel='Density'>
```



```
In [10]: a1=a[['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K', 'Drug']]
```

```
In [11]: sns.heatmap(a1.corr())
```

```
Out[11]: <AxesSubplot:>
```



To Train the Model - Model Building

We are going to train Linear Regression model; We need to split out data into two variables x and y where x is independent variable (input) and y is dependent on x (output). We could ignore address column as it is not required for our model.

```
In [12]: x=a1[['Age']]  
y=a1['Na_to_K']
```

To split my dataset into training and test data

```
In [13]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [14]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[14]: LinearRegression()
```

```
In [15]: print(lr.intercept_)  
18.48275235837274
```


```
In [16]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

```
Out[16]:
```

| | Co-efficient |
|-----|--------------|
| Age | -0.047381 |

```
In [17]: prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[17]: <matplotlib.collections.PathCollection at 0x18eff4cfc70>



```
In [18]: print(lr.score(x_test,y_test))
```

-0.050327243117328724

```
In [19]: from sklearn.linear_model import ElasticNet
en = ElasticNet()
en.fit(x_train,y_train)
```

Out[19]: ElasticNet()

```
In [20]: print(en.coef_)
```

[-0.04547795]

```
In [21]: print(en.intercept_)
```

18.399423447617956

```
In [22]: print(en.predict(x_test))
```

```
[15.21596728 16.80769536 16.9441292 17.12604098 16.62578358 15.35240111
16.76221742 17.12604098 17.62629838 16.4438718 15.62526878 15.67074673
16.12552618 15.89813646 16.12552618 15.48883495 17.39890865 16.85317331
15.71622467 16.9441292 17.30795276 16.76221742 17.35343071 17.39890865
15.48883495 16.26196002 16.48934974 15.62526878 17.67177632 17.03508509
16.4438718 16.80769536 16.62578358 16.30743796 16.98960714 15.30692317
16.53482769 15.07953344 17.35343071 16.76221742 16.21648207 15.26144522
17.35343071 15.07953344 15.53431289 17.39890865 16.62578358 16.4438718
15.12501139 15.76170262 16.26196002 16.26196002 15.67074673 15.57979084
15.85265851 15.26144522 15.35240111 16.9441292 16.17100413 16.08004824]
```

```
In [23]: print(en.score(x_test,y_test))
```

```
-0.04883932126284374
```

```
In [24]: from sklearn import metrics
```

```
In [25]:
```

```
print("Mean Absolytre Error:",metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolytre Error: 5.772271662498251
```