

# Problem Statement

A real estate agent want help to predict the house price for regions in USA.He gave us the dataset to work on to use linear regression model.Create a model that helps him to estimate of what the house would sell for

# Import libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: # To import dataset
df=pd.read_csv('16 Sleep csv')
df
```

Out[2]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Blood Pressure	Heart Rate
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	126/83	77
1	2	Male	28	Doctor	6.2	6	60	8	Normal	125/80	75
2	3	Male	28	Doctor	6.2	6	60	8	Normal	125/80	75
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85
...	...	...	...	...	...	...	...	...	...	...	...
369	370	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	68
370	371	Female	59	Nurse	8.0	9	75	3	Overweight	140/95	68
371	372	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	68
372	373	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	68
373	374	Female	59	Nurse	8.1	9	75	3	Overweight	140/95	68

374 rows × 13 columns

```
In [3]: # To display top 10 rows
df.head(10)
```

Out[3]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Blood Pressure	Heart Rate	Daily Steps
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	126/83	77	4000
1	2	Male	28	Doctor	6.2	6	60	8	Normal	125/80	75	10000
2	3	Male	28	Doctor	6.2	6	60	8	Normal	125/80	75	10000
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85	3000
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85	3000
5	6	Male	28	Software Engineer	5.9	4	30	8	Obese	140/90	85	3000
6	7	Male	29	Teacher	6.3	6	40	7	Obese	140/90	82	3000
7	8	Male	29	Doctor	7.8	7	75	6	Normal	120/80	70	8000
8	9	Male	29	Doctor	7.8	7	75	6	Normal	120/80	70	8000
9	10	Male	29	Doctor	7.8	7	75	6	Normal	120/80	70	8000

# Data Cleaning and Pre-Processing

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 374 entries, 0 to 373
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Person ID                            374 non-null    int64
1   Gender                               374 non-null    object
2   Age                                   374 non-null    int64
3   Occupation                           374 non-null    object
4   Sleep Duration                       374 non-null    float64
5   Quality of Sleep                     374 non-null    int64
6   Physical Activity Level              374 non-null    int64
7   Stress Level                         374 non-null    int64
8   BMI Category                         374 non-null    object
9   Blood Pressure                       374 non-null    object
10  Heart Rate                           374 non-null    int64
11  Daily Steps                          374 non-null    int64
12  Sleep Disorder                       374 non-null    object
dtypes: float64(1), int64(7), object(5)
memory usage: 38.1+ KB
```

In [5]: `df.describe()`

Out[5]:

	Person ID	Age	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	Heart Rate	Daily Steps
<b>count</b>	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000
<b>mean</b>	187.500000	42.184492	7.132086	7.312834	59.171123	5.385027	70.165775	6816.844920
<b>std</b>	108.108742	8.673133	0.795657	1.196956	20.830804	1.774526	4.135676	1617.915679
<b>min</b>	1.000000	27.000000	5.800000	4.000000	30.000000	3.000000	65.000000	3000.000000
<b>25%</b>	94.250000	35.250000	6.400000	6.000000	45.000000	4.000000	68.000000	5600.000000
<b>50%</b>	187.500000	43.000000	7.200000	7.000000	60.000000	5.000000	70.000000	7000.000000
<b>75%</b>	280.750000	50.000000	7.800000	8.000000	75.000000	7.000000	72.000000	8000.000000
<b>max</b>	374.000000	59.000000	8.500000	9.000000	90.000000	8.000000	86.000000	10000.000000

In [6]: `df.columns`

Out[6]: Index(['Person ID', 'Gender', 'Age', 'Occupation', 'Sleep Duration', 'Quality of Sleep', 'Physical Activity Level', 'Stress Level', 'BMI Category', 'Blood Pressure', 'Heart Rate', 'Daily Steps', 'Sleep Disorder'], dtype='object')

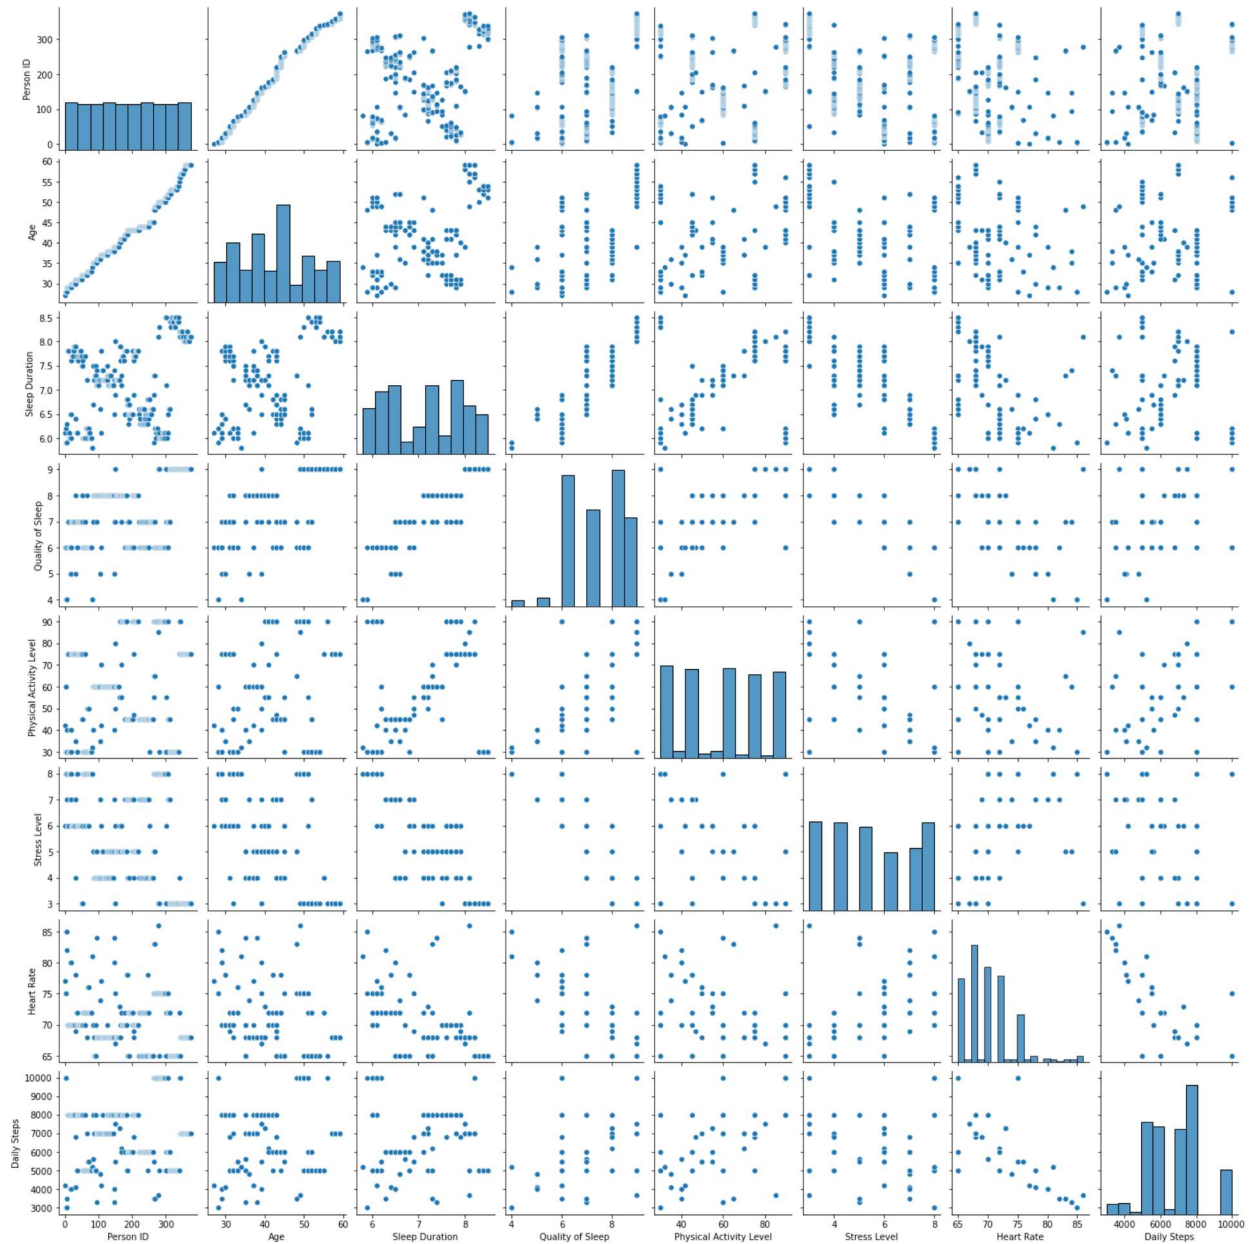
In [7]: `a = df.dropna(axis='columns')`  
`a.columns`

Out[7]: Index(['Person ID', 'Gender', 'Age', 'Occupation', 'Sleep Duration', 'Quality of Sleep', 'Physical Activity Level', 'Stress Level', 'BMI Category', 'Blood Pressure', 'Heart Rate', 'Daily Steps', 'Sleep Disorder'], dtype='object')

## EDA and Visualization

```
In [8]: sns.pairplot(a)
```


```
Out[8]: <seaborn.axisgrid.PairGrid at 0x27864d961f0>
```



In [9]: `sns.distplot(a['Sleep Duration'])`

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)


Out[9]: <AxesSubplot:xlabel='Sleep Duration', ylabel='Density'>



In [10]: `a1=a[['Person ID', 'Age', 'Sleep Duration','Quality of Sleep', 'Physical Activity Level`

In [11]: `sns.heatmap(a1.corr())`

Out[11]: <AxesSubplot:>



## To Train the Model - Model Building

We are going to train Linear Regression model; We need to split out data into two variables x and y where x is the independent variable (Age) and y is the dependent variable (Stress Level). We could train our model using the following code:

```
In [12]: x=a1[['Person ID', 'Age', 'Sleep Duration','Quality of Sleep', 'Physical Activity Level']
y=a1['Stress Level']
```

## To split my dataset into training and test data

```
In [13]: from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [14]: from sklearn.linear_model import LinearRegression

lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[14]: LinearRegression()
```

```
In [15]: print(lr.intercept_)

3.099742684753437e-13
```


```
In [16]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

```
Out[16]:
```

	Co-efficient
<b>Person ID</b>	9.738430e-17
<b>Age</b>	1.930084e-16
<b>Sleep Duration</b>	1.000000e+00
<b>Quality of Sleep</b>	-4.931132e-16
<b>Physical Activity Level</b>	1.385472e-17
<b>Stress Level</b>	-1.777827e-17
<b>Heart Rate</b>	3.729460e-17
<b>Daily Steps</b>	-4.912326e-17

```
In [17]: prediction=lr.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[17]: <matplotlib.collections.PathCollection at 0x2786a393550>
```



```
In [18]: print(lr.score(x_test,y_test))
```

```
1.0
```

## ACCURACY

```
In [19]: from sklearn.linear_model import Ridge,Lasso
```

```
In [20]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_test,y_test)
rr.score(x_train,y_train)
```

```
Out[20]: 0.9901000175495261
```

```
In [21]: rr.score(x_test,y_test)
```

```
Out[21]: 0.9864939832344124
```

```
In [22]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[22]: Lasso(alpha=10)
```

```
In [23]: la.score(x_test,y_test)
```

```
Out[23]: 0.01650048922838787
```

```
In [24]: from sklearn.linear_model import ElasticNet
en = ElasticNet()
en.fit(x_train,y_train)
```

```
Out[24]: ElasticNet()
```

In [25]:

```
print(en.coef_)  
  
[ 0.00096002  0.          0.          0.          0.02505612 -0.  
 -0.08448836 -0.000312  ]
```

In [26]:

```
print(en.intercept_)  
  
13.52071771606479
```

In [27]:

```
print(en.predict(x_test))  
  
[6.93262849 7.19278867 6.70411413 7.02909079 7.82783679 6.59177357  
 6.70123406 6.9047878 6.99163992 6.75802508 6.5341285 6.92974842  
 7.53453571 6.91195795 6.93166846 6.61002732 7.03389091 6.28700756  
 6.7036021 6.59849374 6.68299361 6.86395677 7.53841861 7.82303667  
 6.65797156 6.91342801 6.6954901 7.04445117 7.20142888 6.89230749  
 6.68587368 6.6090673 8.08412555 6.92251821 7.19374869 6.69835399  
 5.87660498 6.93358851 6.70315411 7.56045635 7.20526898 6.99837003  
 7.53457851 7.51297668 7.81535648 7.5422159 7.52881837 7.01469043  
 7.00989031 7.17742829 7.01853053 6.91246799 7.52737703 6.70219409  
 7.23452088 6.9301984 6.94289352 7.72991313 7.81727653 6.69739397  
 8.06065641 7.53933583 7.22350943 6.91579805 7.0175705 6.68203359  
 7.1860685 7.49665628 6.92878839 7.54513877 7.21198914 7.21390919  
 7.53169844 6.92539828 6.90958792 7.5340972 6.88411727 7.53217715  
 7.20334893 6.59081355 7.52493547 6.66187309 7.80095612 6.91054794  
 6.91726811 7.03581095 6.68011354 7.51009661 7.53121713 6.8884674  
 6.50963224 7.03197086 7.50049637 7.21486921 6.5888935 7.56909656  
 6.70264208 6.68875375 7.5259383 7.82207664 7.45006999 7.4803007  
 5.87564496 7.037731 7.03762141 6.8773971 6.58697345 7.82495671  
 7.82687676 6.90670785 7.49329852 7.57389668 7.80863631]
```

In [28]:

```
print(en.score(x_test,y_test))  
  
0.4428043811137372
```

In [29]:

```
from sklearn import metrics  
print("Mean Absolytre Error:",metrics.mean_absolute_error(y_test,prediction))  
print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))  
print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))  
  
Mean Absolytre Error: 6.173233016216623e-14  
Mean Squared Error: 5.134549029614454e-27  
Root Mean Squared Error: 7.165576759490093e-14
```